

62.9 for naive      94.4 for  
naive

Compute and memory to clearly identify the highest contributor. High-level overview of the utilization for compute and memory resource	
Compute (SM) Throughput [%]	70.60
Memory Throughput [%]	84.18
L1/TEX Cache Throughput [%]	88.80
L2 Cache Throughput [%]	8.95
DRAM Throughput [%]	0.78

Compute Throughput Breakdown	
SM: Inst Executed Pipe Lsu [%]	70.60
SM: Mio2rf Writeback Active [%]	40.13
SM: Issue Active [%]	39.28
SM: Inst Executed [%]	39.25
SM: Mio Inst Issued [%]	36.95
SM: Pipe Fma Cycles Active [%]	26.54
SM: Pipe Alu Cycles Active [%]	12.44
SM: Inst Executed Pipe Adu [%]	6.60
SM: Mio Pg Read Cycles Active [%]	3.25
SM: Mio Pg Write Cycles Active [%]	3.25
SM: Pipe Tensor Cycles Active [%]	0.86
SM: Pipe Shared Cycles Active [%]	0.86
SM: Inst Executed Pipe Uniform [%]	0.85
SM: Inst Executed Pipe Cbu Pred On Any [%]	0.05
SM: Inst Executed Pipe Ipa [%]	0
SM: Inst Executed Pipe Fp16 [%]	0
SM: Inst Executed Pipe Tex [%]	0
SM: Pipe Fp64 Cycles Active [%]	0
SM: Inst Executed Pipe Xu [%]	0
IDC: Request Cycles Active [%]	0

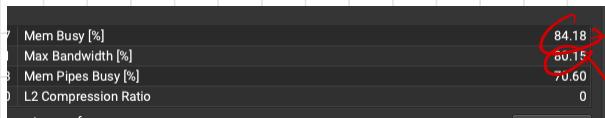
88.80 → slightly less than 89.28 in naive 21  
 up from 68 in naive  
 (due to screen refresh)  
 (but we're using screen now)  
 up from 32% in naive  
 (lots of register file writes/reads)  
 down from 31% in naive (likely due to)

(each of vectorized  
loads)

L1 hit rate

88 → 1.7 b/c of

string



down b/c less strain  
on L1

up b/c  
more strain +  
L1 units constantly  
working

	Instructions	Requests	Waivers	% Peak	Bank Conflicts
Shared Load	41,943,040	41,943,040	50,365,622	76.90	12,418
Shared Load Matrix	0	0	0	0	264,385
Shared Store	2,097,152	2,097,152	2,363,830	3.61	4,216
Shared Store From Global Load	0	0	0	0	0
Shared Atomic	0	0	0	0	0
Other	0	0	0	0	0
Total	4,040,192	44,040,192	62,864,619	80.72	281,021

very good  
arithmetic intensity  
Strm 1x load/lse 20x

very good  
strain usage  
for strain → mrs

L2 Cache										L2 Cache										
	Requests	Sectors	Sections/Rq	% Peak	Hit Rate	Bytes	Throughput	Sectors	Sections/Rq	% Peak	Hit Rate	Bytes	Throughput	Sectors	Sections/Rq	% Peak	Hit Rate	Bytes	Throughput	
L1TEX Store	37768	131072	4	0.19	100	4,194,140,304	7,941,40,3554	65,536	131,072	2	0.12	100	4,194,304	4,873,46,396,57	0	0	0	0	0	
L1TEX Atomic ALU	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
L1TEX Atomic DAS	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
L1TEX Reduction	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
L1TEX Total	2,095,668	8,881,712	4,000	0.98	94.48	248,205,1984	567,743,241,40436	11,975,825	16,226,719	1.35	10.76	97.39	519,255,008	603,562,762,987,70	0	0	0	0	0	0
GCC Total	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
ECC Total	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	
L2 Fabric Total	98,134	344,938	3,511	0.75	92.78	1,101,7356	70,841,88,519,985	35,680	19,798	1.37	0.47	72.38	1,123,7659	1,108,385,698,71	0	0	0	0	0	0
GPU Total	2,035,233	8,455,370	1,97	0.31	94.00	27,579,340	528,386,22,740,29	12,192,003	16,389,549	1.36	0.47	72.38	550,065,366,500	6,177,059,450,75	0	0	0	0	0	0
GPU Total	2,035,233	8,455,370	1,97	0.31	94.00	27,579,340	528,386,22,740,29	12,192,003	16,389,549	1.36	0.47	72.38	550,065,366,500	6,177,059,450,75	0	0	0	0	0	0
L2 Cache Evolution Policies										L2 Cache Evolution Policies										
	First	Hit Rate	Last	Hit Rate	Normal	Hit Rate	Normal	Hit Rate	Last	Hit Rate	Normal	Hit Rate	Normal	Hit Rate	Normal	Hit Rate	Normal	Hit Rate	Normal	
L1TEX Load	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
L1TEX Same	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
L1TEX Atomic	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
L1TEX Total	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
L2 Fabric Total	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
GPU Total	2,565	160	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
Device Memory										Device Memory										
	Sectors	% Peak	Bytes	% Peak	Bytes	Throughput	Sectors	% Peak	Bytes	Throughput	Sectors	% Peak	Bytes	Throughput	Sectors	% Peak	Bytes	Throughput		
Load	262,148	0.78	8,388,736	15,866,52,000,42	0	202,148	0	0	0	0	9,151,478,008,55	0	9,151,478,008,55	0	0	0	0	0	0	
State	0	0	0	0	0	0	262,148	0.78	8,388,736	15,866,52,000,42	0	202,148	0.48	8,388,736	9,151,478,008,55	0	9,151,478,008,55	0		
Total	262,148	0.78	8,388,736	15,866,52,000,42	0	202,148	0	0	0	0	9,151,478,008,55	0	9,151,478,008,55	0	0	0	0	0	0	
Scheduler Statistics										Scheduler Statistics										

in L1tex

sum is now

faster

In native L2 scales we

responsibly

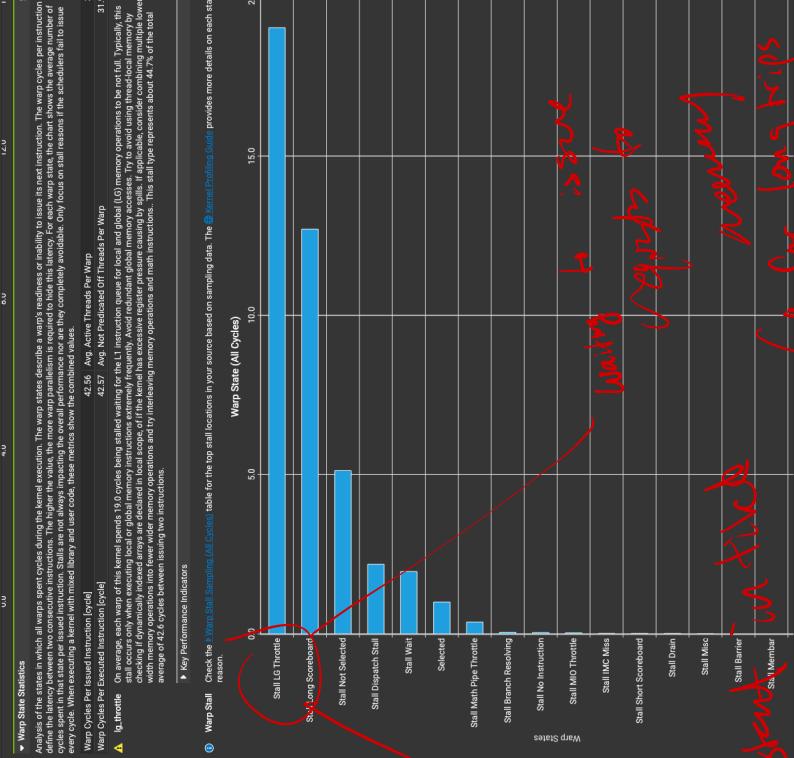
faster is faster and scales given reports

No. divisible

→ sum

(65) nine to

37 in fleet  
due to latency hiding (



Some people  
are born  
to lead

A red ink drawing on graph paper featuring four distinct, abstract, flowing shapes. The top-left shape is a wide, curved loop. The top-right shape is a narrower, more vertical loop. The bottom-right shape is a smaller, rounded loop. The bottom-left shape is a long, thin, horizontal stroke ending in a small circle.